

Principles of Software Construction

Version Control in the Wild

Claire Le Goues

Bogdan Vasilescu




BRANCH WORKFLOWS

<https://www.atlassian.com/git/tutorials/comparing-workflows>


1. Centralized workflow

- Central repository to serve as the single point-of-entry for all changes to the project
- Default development branch is called main
 - all changes are committed into main
 - doesn't require any other branches



Example

John works on his feature




Example

Mary works on her feature



Example


John publishes his feature



Example


John publishes his feature

git push origin main




Example

Mary tries to publish her feature



```
error: failed to push some refs to '/path/to/repo.git'  
hint: Updates were rejected because the tip of your current branch is behind its  
remote counterpart. Merge the remote changes (e.g. 'git pull') before pushing again.  
See the 'Note about fast-forwards' in 'git push --help' for details.
```

Mary tries to publish her feature




Example

Mary rebases on top of John's commit(s)



Mary's Repository




Example

Mary resolves a merge conflict




Example



Example

Mary successfully publishes her feature




2. Git Feature Branch Workflow

- All feature development should take place in a dedicated branch instead of the master branch
- Multiple developers can work on a particular feature without disturbing the main codebase
 - main branch will never contain broken code (enables CI)
 - Enables pull requests (code review)

Example


Mary begins a new feature



```
git checkout -b marys-feature main  
git status  
git add <some-file>  
git commit
```

Example


Mary goes to lunch



```
git push -u origin marys-feature
```


Example

Mary finishes her feature




git push

Example




Example

Mary makes the changes




Example - Merge pull request

Mary publishes her feature




```
git checkout main  
git pull  
git pull origin marys-feature  
git push
```

3. Gitflow Workflow




- Strict branching model designed around the project release
 - Suitable for projects that have a scheduled release cycle
- Branches have specific roles and interactions
- Uses two branches
 - main stores the official release history; tag all commits in the main branch with a version number
 - develop serves as an integration branch for features

GitFlow feature branches (from develop)




GitFlow release branches (eventually into master)



no new features after this point—only bug fixes, docs, and other release tasks

GitFlow hotfix branches

used to quickly patch production releases



Aside: Semantic Versioning

Semantic Versioning

Given a version number MAJOR.MINOR.PATCH, increment the:

1. MAJOR version when you make incompatible API changes,
2. MINOR version when you add functionality in a backwards compatible manner, and
3. PATCH version when you make backwards compatible bug fixes.

Code status	Stage	Rule	Example version
First release	New product	Start with 1.0.0	1.0.0
Backward compatible bug fixes	Patch release	Increment the third digit	1.0.1
Backward compatible new features	Minor release	Increment the middle digit and reset last digit to zero	1.1.0
Changes that break backward compatibility	Major release	Increment the first digit and reset middle and last digits to zero	2.0.0

Summary

- Version control has many advantages
 - History, traceability, versioning
 - Collaborative and parallel development
- Collaboration with branches
 - Different workflows
- From local to central to distributed version control


DEVELOPMENT AT SCALE

Releasing at scale in industry

- Facebook:
<https://atscaleconference.com/videos/rapid-release-at-massive-scale/>
- Google:
<https://www.slideshare.net/JohnMicco1/2016-0425-continuous-integration-at-google-scal>
<https://testing.googleblog.com/2011/06/testing-at-speed-and-scale-of-google.html>
- Why Google Stores Billions of Lines of Code in a Single Repository:
<https://www.youtube.com/watch?v=W71BTkUbdqE>
- F8 2015 - Big Code: Developer Infrastructure at Facebook's Scale:
<https://www.youtube.com/watch?v=X0VH78ye4yY>

Pre-2017 release management model at Facebook

Diff lifecycle: local testing



Tools/xctool/xctool/xctool/Version.m

```
1 #import "Version.h"
2
3 NSString * const XCToolVersionString = @".0.2.1";
```

View Options ▾

```
1 #import "Version.h"
2
3 NSString * const XCToolVersionString = @".0.2.2";
```

```
PASS ExampleTest (0.050s)
.
OK (1 test, 4 assertions)
OK
(1 tests, 4 assertions, 0 incomplete, 0 failures)
```

Test and lint locally


Diff lifecycle: CI testing (data center)



	Facebook	Messenger	Groups	...
arm	✓	✓	✓	✓
x86	✓	✓	✓	✓
...	✓	✓	✓	✓


App and Build
Configuration Matrix

Diff lifecycle: diff ends up on main branch




Release every two weeks

www.facebook.com



Quasi-continuous push from master (1,000+ devs, 1,000 diffs/day); 10 pushes/day



<https://samritchie.wordpress.com/2013/10/16/build-server-traffic-lights/>



<https://www.softwire.com/blog/2013/09/26/continuous-integration-traffic-lights-revamp/index.html>

Status

Build Pipeline




Release Pipeline


Dev	Test	Prod
deployment succeeded	deployment succeeded	deployment succeeded
NuGet 0.6.0	NuGet 0.6.0	NuGet 0.4.0

<https://blog.devops4me.com/status-badges-in-azure-devops-pipelines/>

Signals of PR quality



Result: Build status+code coverage badges indicate *more tests in PRs*




Aside: Key idea – fast to deploy, slow to release

Dark launches at Instagram

- **Early:** Integrate as soon as possible. Find bugs early. Code can run in production about 6 months before being publicly announced (“dark launch”).
- **Often:** Reduce friction. Try things out. See what works. Push small changes just to gather metrics, feasibility testing. Large changes just slow down the team. Do dark launches, to see what performance is in production, can scale up and down. *"Shadow infrastructure" is too expensive, just do in production.*
- **Incremental:** Deploy in increments. Contain risk. Pinpoint issues.

Aside: Feature Flags

Typical way to implement a dark launch.



<http://swreflections.blogspot.com/2014/08/feature-toggles-are-one-of-worst-kinds.html>

<http://martinfowler.com/bliki/FeatureToggle.html>


Issues with feature flags

Feature flags are “technical debt”


Example: \$400 million financial services company went bankrupt in 45 minutes.

<http://dougseven.com/2014/04/17/knightmare-a-devops-cautionary-tale/>


Diff lifecycle: in production



What's in a weekly branch cut? (The limits of branches)



Post-2017 release management model at Facebook



Google: similar story. YUGE code base

Google repository statistics


As of Jan 2015

Total number of files*	1 billion
Number of source files	9 million
Lines of code	2 billion
Depth of history	35 million commits
Size of content	86 terabytes
Commits per workday	45 thousand

*The total number of files includes source files copied into release branches, files that are deleted at the latest revision, configuration files, documentation, and supporting data files.

Exponential growth

Millions of changes committed (cumulative)



Google Speed and Scale

- >30,000 developers in 40+ offices
- 13,000+ projects under active development
- 30k submissions per day (1 every 3 seconds)

- All builds from source
- 30+ sustained code changes per minute with 90+ peaks
- 50% of code changes monthly
- 150+ million test cases / day, > 150 years of test / day
- Supports continuous deployment for all Google teams!

Google Confidential and Proprietary

Google code base vs Linux kernel code base

Some perspective

Linux kernel

- 15 million lines of code in 40 thousand files (total)

Google repository

- 15 million lines of code in 250 thousand files *changed per week, by humans*
- 2 billion lines of code, in 9 million source files (total)

How do they do it?

1. Lots of (automated) testing

Google workflow



- All code is reviewed before commit (by humans and automated tooling)
- Each directory has a set of owners who must approve the change to their area of the repository
- Tests and automated checks are performed before and after commit
- Auto-rollback of a commit may occur in the case of widespread breakage

2. Lots of automation

Additional tooling support

Critique	Code review
CodeSearch*	Code browsing, exploration, understanding, and archeology
Tricorder**	Static analysis of code surfaced in Critique, CodeSearch
Presubmits	Customizable checks, testing, can block commit
TAP	Comprehensive testing before and after commit, auto-rollback
Rosie	Large-scale change distribution and management

* See "How Developers Search for Code: A Case Study", In European Software Engineering Conference and the ACM SIGSOFT Symposium on the Foundations of Software Engineering, 2015

** See "Tricorder: Building a program analysis ecosystem", In International Conference on Software Engineering (ICSE), 2015


3. Smarter tooling

- Build system
- Version control
- ...

3a. Build system

Google Standard Continuous Build System


- Triggers builds in continuous cycle
- Cycle time = longest build + test cycle
- Tests many changes together
- Which change broke the build?



Google Confidential and Proprietary

Google Google Continuous Build System

- Triggers tests on every change
- Uses fine-grained dependencies
- Change 2 broke test 1



Google Confidential and Proprietary

Google Continuous Integration Display

Google Benefits

- Identifies failures sooner
- Identifies culprit change precisely
 - Avoids divide-and-conquer and tribal knowledge
- Lower compute costs using fine grained dependencies
- Keeps the build green by reducing time to fix breaks
- Accepted enthusiastically by product teams
- Enables teams to ship with fast iteration times
 - Supports submit-to-production times of less than 36 hours for some projects

Google Confidential and Proprietary

Google Costs

- Requires enormous investment in compute resources (it helps to be at Google) grows in proportion to:
 - Submission rate
 - Average build + test time
 - Variants (debug, opt, valgrind, etc.)
 - Increasing dependencies on core libraries
 - Branches
- Requires updating dependencies on each change
 - Takes time to update - delays start of testing

Which tests to run?

GMAIL


Test Target:

```
name: //depot/gmail_client_tests  
name: //depot/gmail_server_tests
```


BUZZ

Test targets:


```
name: //depot/buzz_server_tests  
name: //depot/buzz_client_tests
```




Scenario 1: a change modifies common_collections_util



Scenario 1: a change modifies common_collections_util




Scenario 1: a change modifies common_collections_util




Scenario 1: a change modifies common_collections_util

All tests are affected! Both Gmail and Buzz projects need to be updated




Scenario 2: a change modifies the youtube_client



Scenario 2: a change modifies the youtube_client

Only buzz_client_tests are run and only Buzz project needs to be updated.




3b. Version control

- Problem: even git can get slow at Facebook scale
 - 1M+ source control commands run per day
 - 100K+ commits per week



3b. Version control

- Solution: redesign version control




3b. Version control

- Solution: redesign version control
 - Query build system's file monitor, Watchman, to see which files have changed

3b. Version control

- Solution: redesign version control
 - Query build system's file monitor, Watchman, to see which files have changed → **5x faster “status” command**



3b. Version control


- Solution: redesign version control
 - Sparse checkouts??? (remember, git is a distributed VCS)

3b. Version control


- Solution: redesign version control
 - Sparse checkouts:
 - Change the clone and pull commands to download only the commit metadata, while omitting all file changes (the bulk of the download)
 - When a user performs an operation that needs the contents of files (such as checkout), download the file contents on demand using existing memcache infrastructure

3b. Version control

- Solution: redesign version control
 - Sparse checkouts → **10x faster clones and pulls**
 - Change the clone and pull commands to download only the commit metadata, while omitting all file changes (the bulk of the download)
 - When a user performs an operation that needs the contents of files (such as checking out a branch), the system can then fetch the files from the repository and use existing memcache infrastructure to store them.



4. Monolithic repository




Monolithic repository – no major use of branches for development


Trunk-based development

Combined with a centralized repository, this defines the monolithic model


- Piper users work at “head”, a consistent view of the codebase
- All changes are made to the repository in a single, serial ordering
- There is no significant use of branching for development
- Release branches are cut from a specific revision of the repository



Did it work? Yes. Sustained productivity at Facebook



Lines Committed Per Developer Per Day



Growth of the size of the Android and iOS dev teams

to be continued ...

MONOREPO VS MANY REPOS

A recent history of code organization

- A single team with a monolithic application in a single repository
- ...
- Multiple teams with many separate applications in many separate repositories
- Multiple teams with many ~~separate applications~~ microservices in many separate repositories
- A single team with many microservices in many repositories
- ...
- Many teams with many applications in one big **Monorepo**

What is a Monolithic Repository (monorepo)?

A **single** version control repository containing multiple

- ▶ projects
- ▶ applications
- ▶ libraries,

often using a common build system.

Monorepos in industry

Google (computer science version)

TRUSTED INSIGHTS FOR COMPUTING'S LEADING PROFESSIONALS ACM.org | Join ACM | About Communications | ACM Resources | Alerts & Feeds    

SIGN IN

COMMUNICATIONS OF THE ACM

HOME CURRENT ISSUE NEWS BLOGS OPINION RESEARCH PRACTICE CAREERS ARCHIVE VIDEOS

Home / Magazine Archive / July 2016 (Vol. 59, No. 7) / Why Google Stores Billions of Lines of Code in a Single... / Full Text

CONTRIBUTED ARTICLES

Why Google Stores Billions of Lines of Code in a Single Repository

By Rachel Potvin, Josh Levenberg
Communications of the ACM, Vol. 59 No. 7, Pages 78-87
10.1145/2854146
Comments (3)

VIEW AS:       SHARE:      



Early Google employees decided to work with a shared codebase managed through a centralized source control system. This approach has served Google well for more than 16 years, and today the vast majority of Google's software assets continues to be stored in a single, shared repository. Meanwhile, the number of Google software developers has steadily increased, and the size of the Google codebase has grown exponentially (see Figure 1). As a result, the technology used to host the codebase has also evolved significantly.

[Back to Top](#)

SIGN IN for Full Access

User Name

Password

[» Forgot Password?](#)

[» Create an ACM Web Account](#)

SIGN IN

ARTICLE CONTENTS:

Introduction
Key Insights
Google-Scale
Background
Analysis

Monorepos in industry

Scaling Mercurial at Facebook

The screenshot shows a blog post on the Facebook Code website. The header includes a navigation bar with categories like Open Source, Platforms, Infrastructure Systems, Hardware Infrastructure, Video & VR, and Artificial Intelligence. The main content is titled "Scaling Mercurial at Facebook" and is dated January 2014. It features two authors: Durham Goode and Siddharth P Agarwal. The post discusses the challenges of managing a massive codebase and how Facebook's practices have improved their source control. Below the main content, there are sections on choosing a source control system and scaling memcached. A sidebar on the right is titled "Recommended" and lists other posts from the Facebook Code blog.

Scaling Mercurial at Facebook

7 January 2014 INFRA · OPEN SOURCE · PERFORMANCE · OPTIMIZATION

Scaling Mercurial at Facebook

Durham Goode Siddharth P Agarwal

With thousands of commits a week across hundreds of thousands of files, Facebook's main source repository is enormous—many times larger than even the Linux kernel, which checked in at 17 million lines of code and 44,000 files in 2013. Given our size and complexity—and Facebook's practice of shipping code twice a day—improving our source control is one way we help our engineers move fast.

Choosing a source control system

Two years ago, as we saw our repository continue to grow at a staggering rate, we sat down and extrapolated our growth forward a few years. Based on those projections, it appeared likely that our then-current technology, a Subversion server with a Git mirror, would become a productivity bottleneck very soon. We looked at the available options and found none that were both fast and easy to use at scale.

Our code base has grown organically and its internal dependencies are very complex. We could have spent a lot of time making it more modular in a way that would be friendly to a source control tool, but there are a number of benefits to using a single repository. Even at our current scale, we often make large changes throughout our code base, and having a single repository is useful for continuous

Recommended

Scaling memcached at Facebook

Flashcache at Facebook: From 2010 to 2013 and beyond

Monorepos in industry

Microsoft claim the largest git repo on the planet

The screenshot shows a blog post on the Microsoft Developer Tools Blogs. The post is titled "The largest Git repo on the planet" and was published on 05/24/2017 by Brian Harry MS. It has 59 comments and a rating of 4.5 stars. The post discusses the development of Git Virtual File System (GVFS) to handle extremely large projects and teams. It mentions that GVFS enables Git to scale to very large repos by virtualizing both the .git folder and the working directory, allowing it to dynamically download only the portions needed. The post also notes that a lot has happened since the initial write-up, including expanded open sourcing and improvements to how it works for Microsoft, partners, and customers.

Brian Harry's blog
Everything you want to know about Visual Studio ALM and Farming

The largest Git repo on the planet
05/24/2017 by Brian Harry MS // 59 Comments

Share 2.2k | 3711 | 1230

It's been 3 months since I first wrote about our efforts to scale Git to extremely large projects and teams with an effort we called "Git Virtual File System". As a reminder, GVFS, together with a set of enhancements to Git, enables Git to scale to VERY large repos by virtualizing both the .git folder and the working directory. Rather than download the entire repo and checkout all the files, it dynamically downloads only the portions you need based on what you use.

A lot has happened and I wanted to give you an update. Three months ago, GVFS was still a dream. I don't mean it didn't exist – we had a concrete implementation, but rather, it was unproven. We had validated on some big repos but we hadn't rolled it out to any meaningful number of engineers so we had only conviction that it was going to work. Now we have proof.

Today, I want to share our results. In addition, we're announcing the next steps in our GVFS journey for customers, including expanded open sourcing to start taking contributions and improving how it works for us at Microsoft, as well as for partners and customers.

Windows is live on Git

Over the past 3 months, we have largely completed the rollout of Git/GVFS to the Windows team at Microsoft.

As a refresher, the Windows code base is approximately 3.5M files and, when checked in to a Git repo, results in a repo of about 300GB.

Visual Studio

Download Visual Studio | Download TFS | Visual Studio Team Services

Search

Search MSDN with Bing | Search this blog | Search all blogs

Subscribe Blog via Email

Subscribe to this blog and receive notifications of new posts by email.

Email Address: Subscribe! Unsubscribe

institute for SOFTWARE RESEARCH

Monorepos in open-source

foresquare public monorepo

The screenshot shows the GitHub repository page for `foursquare / fsqio`. The repository is described as "A monorepo that holds all of Foursquare's opensource projects". It has 538 commits, 1 branch, 2 releases, 16 contributors, and is licensed under Apache-2.0. The repository was last updated on 1 Aug. The commit history lists various contributors and their changes, including updates to Travis config, Dockerfiles, and global variables.

Commit	Message	Date
mateor	committed with mateor Upgrade Fsq.io Travis config to use mongodb3.0+ (#780)	Latest commit 494b379 on 1 Aug
3rdparty	Update the testinfra deployed file (#748)	3 months ago
build-support	Monolithic Ivy resolve commit (#530)	3 months ago
scripts/fsqio	Add a check for the current file before deleting (#709)	3 months ago
src	Add installation instructions to pom	3 months ago
test	Spindle: Make ThriftParserTest actually depend on its input (#735)	3 months ago
.dockerrignore	Update fsqio/fsqio Dockerfile and add one for fsqio/twofishes	2 years ago
.gitignore	Update upkeep to no longer clobber global variables	10 months ago
.travis.yml	Upgrade Fsq.io Travis config to use mongodb3.0+ (#780)	3 months ago
BUILD.opensource	Monolithic Ivy resolve commit (#530)	3 months ago
BUILD.tools	Drop a BUILD.tools in Fsq.io.	8 months ago
CLA.md	Move deployed files to consolidated directory.	2 years ago
CONTRIBUTING.md	Part 2: CONTRIBUTING.md	2 years ago

Monorepos in open-source

The Symfony monorepo

43 projects, **25 000** commits, and **400 000** LOC

<https://github.com/symfony/symfony>


Bridge/
5 sub-projects

Bundle/
5 sub-projects


Component/
33 independent sub-projects like Asset, Cache,
CssSelector, Finder, Form, HttpKernel, Ldap,
Routing, Security, Serializer, Templating,
Translation, Yaml, ...

Common build system


Bazel from Google



Buck from Facebook



Pants from Twitter



Some advantages of monorepos

High Discoverability For Developers

- ▶ Developers can read and explore the whole codebase
- ▶ grep, IDEs and other tools can search the whole codebase
- ▶ IDEs can offer auto-completion for the whole codebase
- ▶ Code Browsers can links between all artifacts in the codebase

Code-Reuse is cheap

Almost zero cost in introducing a new library

- ▶ Extract library code into a new directory/component
- ▶ Use library in other components
- ▶ Profit!

Refactorings in one commit

Allow large scale refactorings with one single,
atomic, history-preserving commit

- ▶ Extract Library/Component
- ▶ Rename Functions/Methods/Components
- ▶ Housekeeping (phpcs-fixer, Namespacing, ...)

Another refactoring example

- Make large backward incompatible changes easily... especially if they span different parts of the project
- For example, old APIs can be removed with confidence
 - Change an API endpoint code **and** all its usages in **all** projects in **one** pull request

Some more advantages

- Easy continuous integration and code review for changes spanning several projects
- (Internal) dependency management is a non-issue
- Less context switching for developers
- Code more reusable in other contexts
- Access control is easy

Some downsides

- Require collective responsibility for team and developers
- Require trunk-based development
 - Feature toggles are technical debt (recall financial services example)
- Force you to have only one version of everything
- Scalability requirements for the repository
- Can be hard to deal with updates around things like security issues
- Build and test bloat without very smart build system
- Slow VCS without very smart system
- Permissions?

Summary

- Version control has many advantages
 - History, traceability, versioning
 - Collaborative and parallel development
- Collaboration with branches
 - Different workflows
- From local to central to distributed version control

Summary

- Configuration management
 - Treat infrastructure as code
 - Git is powerful
- Release management: versioning, branching, ...
- Software development at scale requires a lot of infrastructure
 - Version control, build managers, testing, continuous integration, deployment, ...
- It's hard to scale development
 - Move towards heavy automation (DevOps)
- Continuous deployment increasingly common
- Opportunities from quick release, testing in production, quick rollback